



Inspire...Educate...Transform.

Applying ML to Big Data using
Hadoop and Spark Ecosystem

Spark and towards Spark ML

Dr. Manoj Duse

Aug 5, 2018

Agenda

- Spark
- RDDs, Partitions
- Actions, Transformers
- Managing partitions
- Spark ML Overview



Another Parallelization
platform: SPARK

- SPARK is not replacement for HADOOP
- SPARK on YARN
- SPARK and MR can coexist
- SPARK for real-time, MR for batch

What is Spark?

Fast and Expressive Cluster Computing System
Compatible with Apache Hadoop

Up to **10x** faster on disk,
100x in memory

Efficient

- General execution graphs
- In-memory storage


2-5x less code

Usable

- Rich APIs in Java, Scala, Python
- Interactive shell



SPARK ecosystem



Spark Core:

Responsible for:

- ✓ Memory Management and fault recovery
- ✓ Supports/implements key concepts of RDDs and Actions
- ✓ Scheduling, Monitoring, Distributing jobs on cluster [via YARN]



Key Concepts

Write programs in terms of transformations
on distributed datasets

Resilient Distributed Datasets

- Collections of objects spread across a cluster, stored in RAM or on Disk
- Built through parallel transformations
- Automatically rebuilt on failure

Operations

- Transformations (e.g. map, filter, groupBy)
- Actions (e.g. count, collect, save)



Spark Terminology

Driver program	The process running the main() function of the application and creating the SparkContext
----------------	--

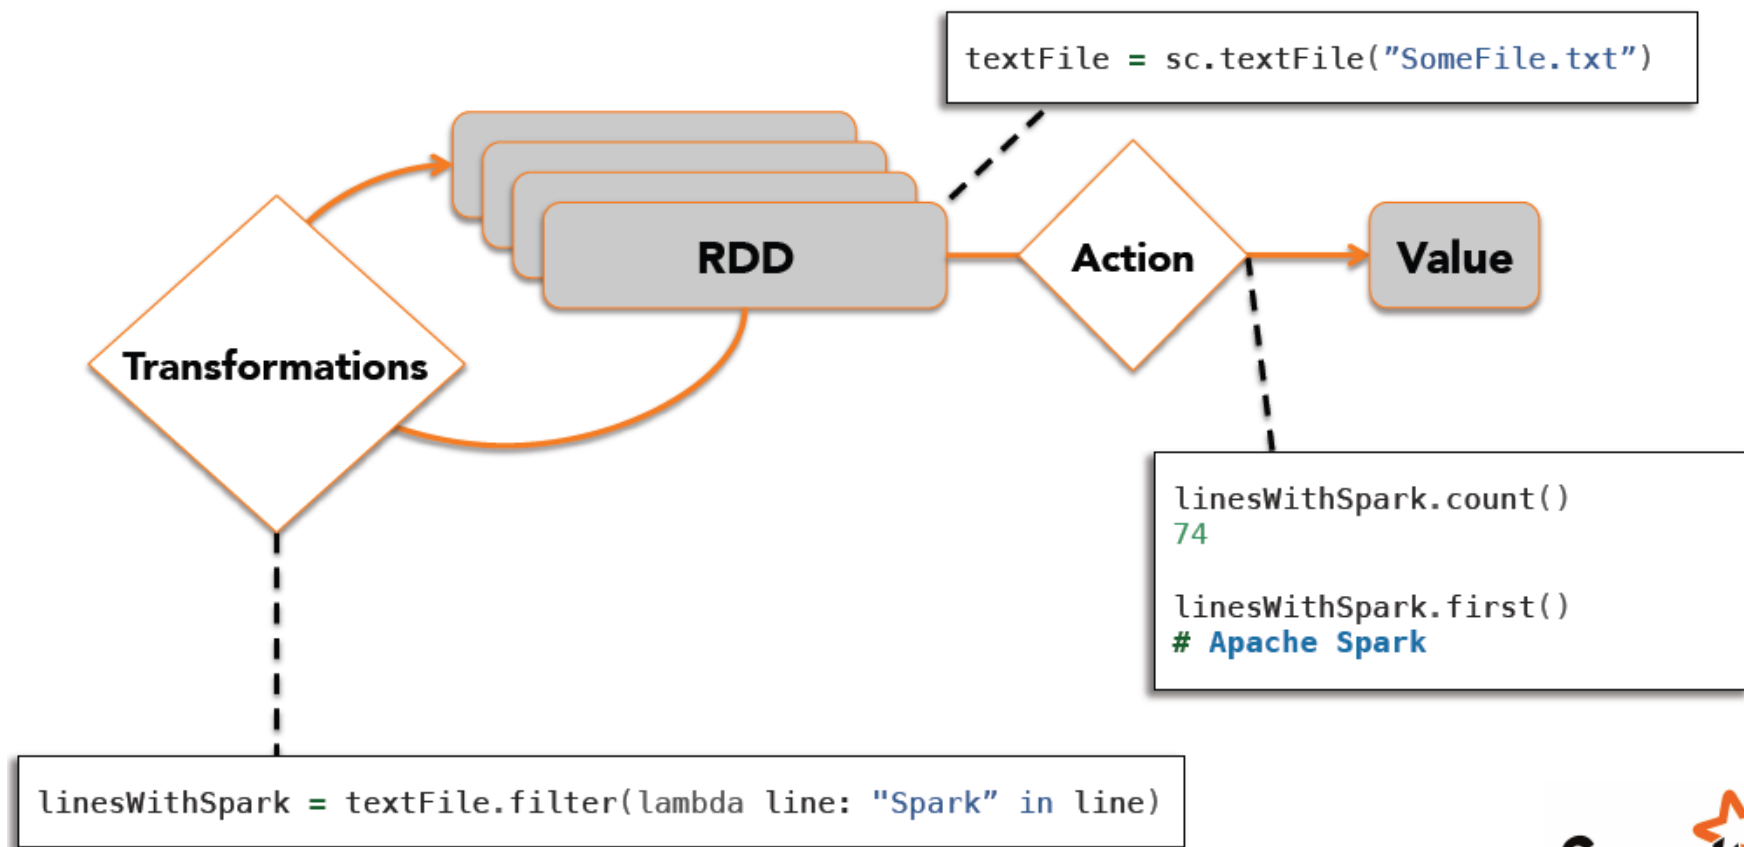
Executor	A process launched for an application on a worker node, that runs tasks and keeps data in memory or disk storage across them. Each application has its own executors.
----------	--

SparkContext works with RM

Programming languages supported by Spark include:

- Java
- Python [PySpark]
- Scala
- SQL
- R [SparkR]

Working With RDDs



RDDs are immutable

Transformation on one RDD results into a new RDD

Spark Example: Log Mining

- Load error messages from a log into memory, then interactively search for various patterns

```
lines = sc.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()


cachedMsgs.filter(_.contains("database")).count
cachedMsgs.filter(_.contains("memory")).count
...
```

Base RDD

Transformed RDD

Cached RDD

Parallel operation



You should cache RDDs you work with when you want to execute two or more actions on it for a better performance

Lazy Initialization

- Populating of blocks into memory deferred until action is invoked.
- RDD created but with no data
- Simply put, an action triggers actual evaluation of the RDD .
- Only actions can materialize the entire processing pipeline with real data.

DAG, Lineage Graph [resilient]

Examples of Transformations

- map
- filter
- flatMap
- groupByKey
- sortByKey

Examples of Actions

- Count
- Top(k)

Broadcast variable and Accumulator

- Broadcast variable is a read-only variable
 - made available from the driver program that runs the SparkContext object to the nodes that will execute the computation.
 - useful in applications that need to make the same [typically reference data] available to the worker nodes in an efficient manner, such as machine learning algorithms.
 - one time thing : distributed to the workers only once
-
- An accumulator is also a variable that is broadcasted to the worker nodes.
 - The key difference between a broadcast variable and an accumulator is that while the broadcast variable is read-only, the accumulator can be added to.

Printing contents of a RDD

- `myRDD.collect().foreach(println)`
- Very useful for debugging

Parallelizing Data

How many partitions my RDD is split into?

```
myRDD.partitions.size
```

How to enforce “degree of parallelism”

```
myRDD = sc.parallelize(1 to 500, 5)
```

```
myRDD.partitions.size
```

```
res27: Int = 5
```

repartition vs coalesce

repartition : will shuffle the original partitions and repartition them

coalesce : will just combine original partitions to the new number of partitions.

Shuffling could be very costly,
If all you want is to reduce the no of partitions: use coalesce

RDD → DataFrames → DataSets

- **Spark Dataframe APIs –**

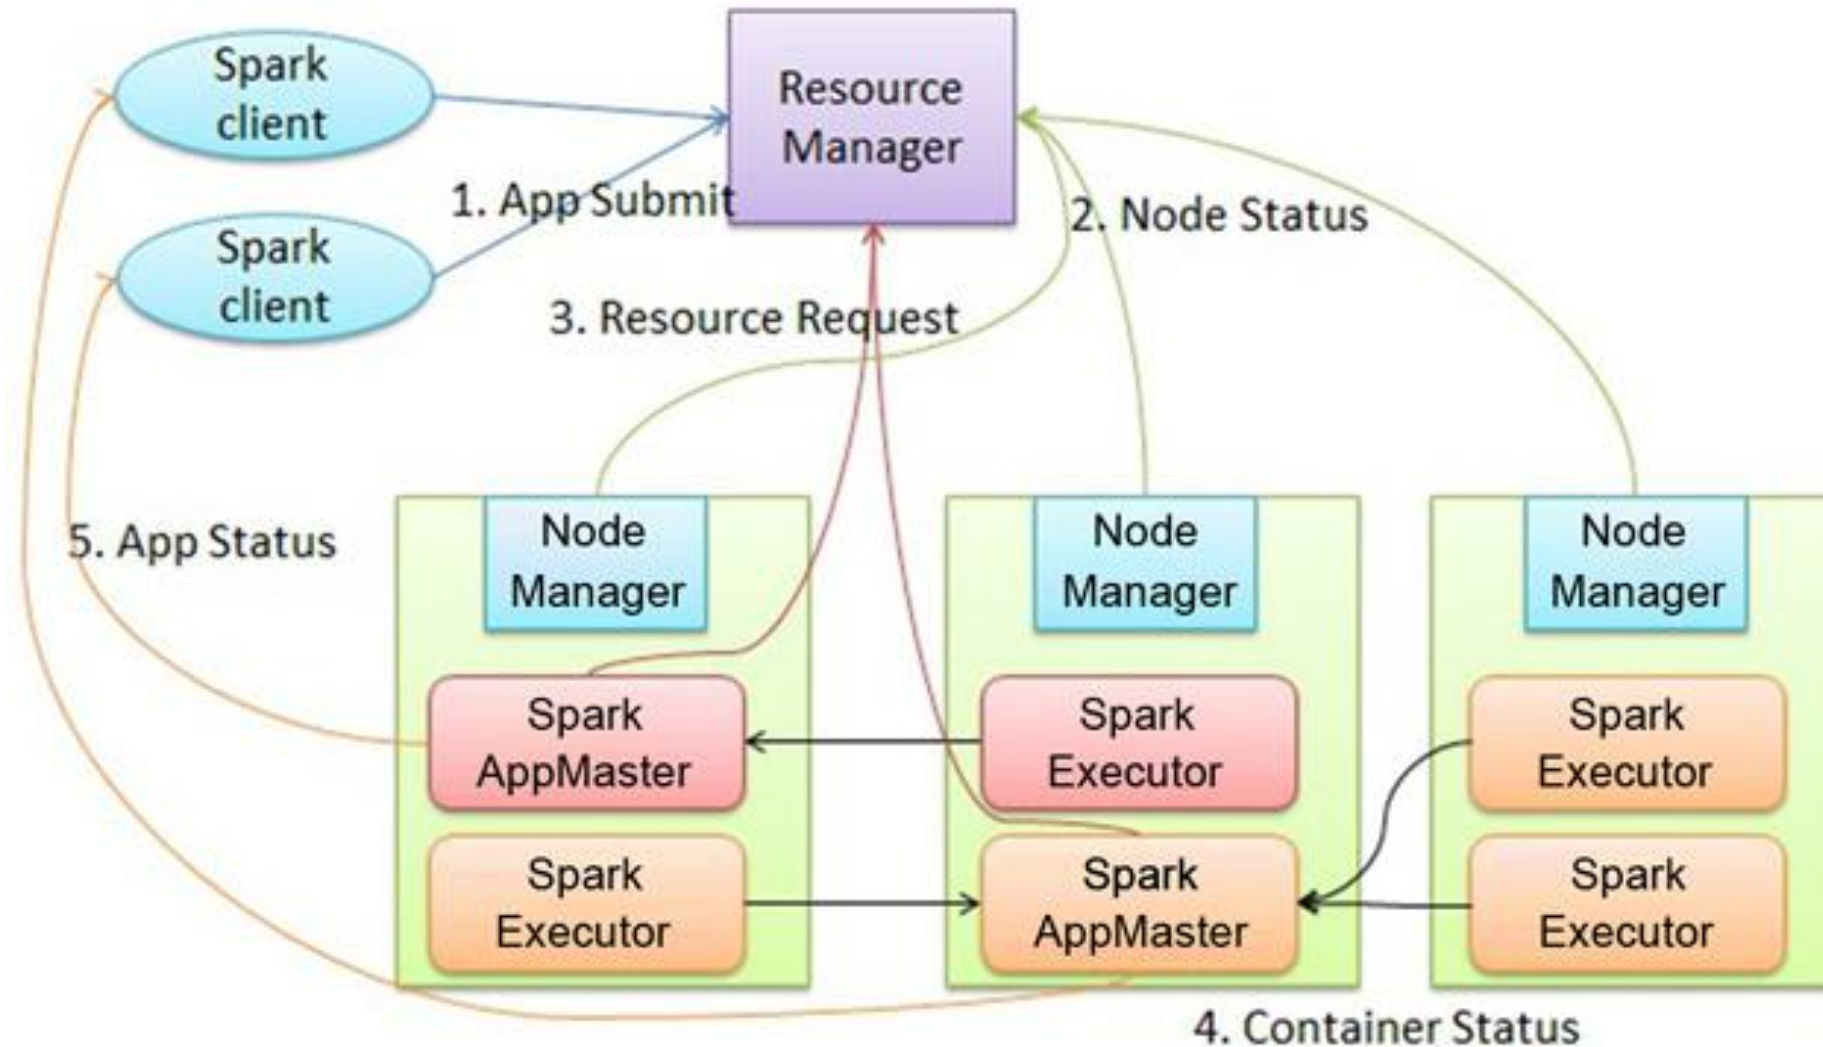
- Unlike an RDD, data organized into named columns.
- For example a table in a relational database.
- Allows developers to impose a structure onto a RDD

- **Spark Dataset APIs –**

- Datasets in Apache Spark are an extension of DataFrame API which provides type-safe [compile time], object-oriented programming interface.
- One can seamlessly move between DataFrame or Dataset and RDDs by simple API method calls like .rdd or .toDF
- DataFrames and Datasets are built on top of RDDs.

- **RDD** – The RDD APIs have been on Spark since the 1.0 release.
- **DataFrames** – Spark introduced DataFrames in Spark 1.3 release.
- **DataSet** – Spark introduced Dataset in Spark 1.6 release.

Relating back to YARN



SPARK vs MR

- Ease of use
- Developer productivity
- Speed
- Ecosystem: Spark R, Spark MLlib, PySpark, SparkSQL
- Lot of apache projects also moving to support/leverage Spark

Spark ML

At a high level, it provides tools such as:

- ML Algorithms: common learning algorithms such as classification, regression, clustering, and collaborative filtering
- Featurization: feature extraction, transformation, dimensionality reduction, and selection
- Pipelines: tools for constructing, evaluating, and tuning ML Pipelines
- Persistence: saving and load algorithms, models, and Pipelines
- Utilities: linear algebra, statistics, data handling, etc.

ML Lib vs ML

ML spark.ml

- New
- DataFrames
- Pipelines

ML Lib spark.mllib

- Old
- RDDs
- But more features but ML catching up
- In maintenance mode; no new functionality will be added

What is the future direction...

- After reaching feature parity (roughly estimated for Spark 2.3), the RDD-based API will be deprecated.
- The RDD-based API is expected to be removed in Spark 3.0.

- **Estimator**: An Estimator is an algorithm which can be fit on a DataFrame to produce a Transformer.

E.g., a learning algorithm is an Estimator which trains on a DataFrame and produces a model.

- **Transformer**: A Transformer is an algorithm which can transform one DataFrame into another DataFrame.

E.g., an ML model is a Transformer

It transforms a "DataFrame with features" → into a "DataFrame with predictions".

- **Pipeline**: A Pipeline chains multiple Transformers and Estimators together to specify an ML workflow.

Pipeline

- A **Estimator implements a method fit()**, which accepts a DataFrame and produces a Model, which is a Transformer.

For example, a learning algorithm such as LogisticRegression is an Estimator

Calling fit() trains a LogisticRegressionModel, which is a Model and hence a Transformer.

- Pipeline, which consists of a sequence of **PipelineStages**

Estimators and Transformers to be run in a specific order

To get a rough idea :

```
val lr = new LogisticRegression() .setMaxIter(10) .setRegParam(0.001)
```

```
val pipeline = new Pipeline() .setStages(Array(tokenizer, hashingTF, lr))
```

```
// Fit the pipeline to training documents.
```

```
val model = pipeline.fit(training)
```

```
// Now we can optionally save the fitted pipeline to disk
```

```
model.write.overwrite().save("/tmp/spark-logistic-regression-model")
```

```
// Make predictions on test documents.
```

```
model.transform(test)
```

So what all we have learnt ?

Web: <http://www.insofe.edu.in>

Facebook: <https://www.facebook.com/insofe>

Twitter: <https://twitter.com/Insofeedu>

YouTube: <http://www.youtube.com/InsofeVideos>

SlideShare: <http://www.slideshare.net/INSOFE>

LinkedIn: <http://www.linkedin.com/company/international-school-of-engineering>